

Baby42

A simple RISC processor

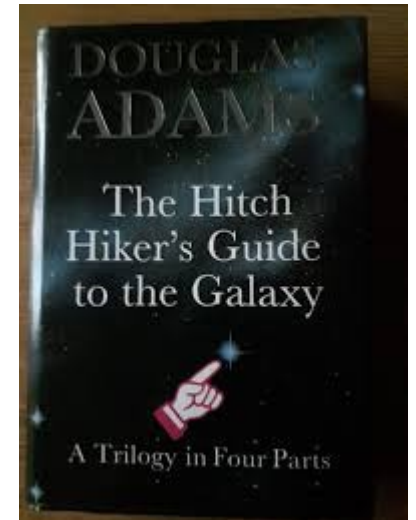
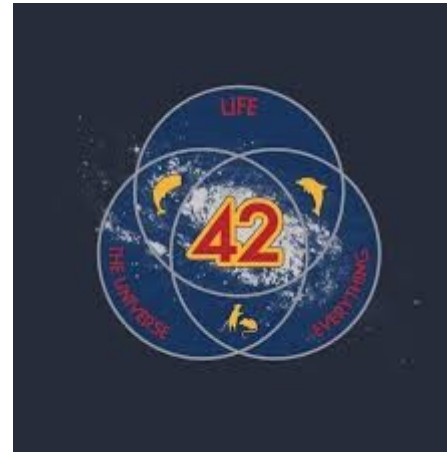
Jecel Mattos de Assumpção Jr
August 2022

name



- The Manchester Baby (Small-Scale Experimental Machine – SSEM) ran a stored memory program in June 1948

- 4 bytes (32 bits) of data with 2 bytes of instruction
- And:



state

General purpose
registers

Program Counter

	31	24	23	16	15	8	7	0
r15								
r14								
r13								
r12								
r11								
r10								
r9								
r8								
r7								
r6								
r5								
r4								
r3								
r2								
r1								
r0								

31	24	23	16	15	8	7	0

17 words = 68 bytes =
136 hexadecimal digits
+ external memory

Manchester Baby: 3 registers
+ 32 words of memory of 32
bits each

RISC style

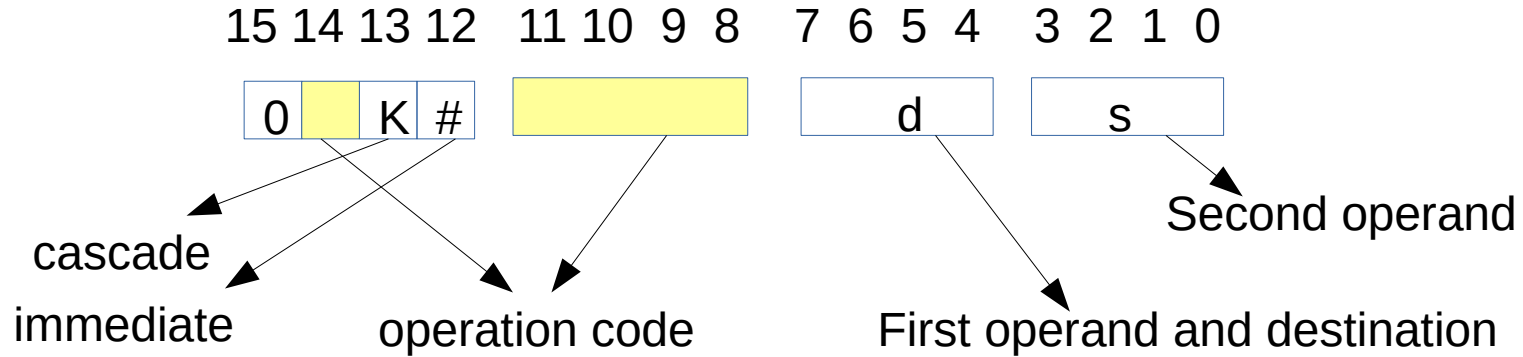
Two instruction formats:

- Data manipulation instructions only operate on registers, or a register and an immediate value present in the instruction itself
- Data storage transfer a register to/from memory (jump and link has the same format, though it should be considered a third kind of instruction)

RISC: how many instructions?

- Counting all variations in assembly syntax we have 106 standard instructions and 12 optional ones (multiplication)
- Not counting the variations we have 30 standard instructions and 3 optional ones
- The control unit deals with only five instruction types:
 - 1)ALU (Math, Logic and Multiply)
 - 2)Comparison
 - 3)Load
 - 4)Store
 - 5)Jump and Link

Data Manipulation Instructions



30 of the 32 possible op codes are used, but due to # and K nearly all instructions have 4 variations.

The assembly language syntax is C-like, with an addition being indicated by

`rD += rS`

- Immediate values

"immediate" changes the second operand to be a number from -6 to 7 instead of a register. The value -7 (9 in hex) indicates that the actual operand is the next 32 bits and -8 (8 in hex) indicates 16 bits

In assembly, a # character is placed before the second operand to indicate that it is an immediate. Otherwise the numbers 0 to 15 indicate a register.

K - Cascade

"cascade" changes the destination to be the first operand of the following instruction. A sequence like

```
r3 <<= #2 ; multiply by four  
r3 += r1
```

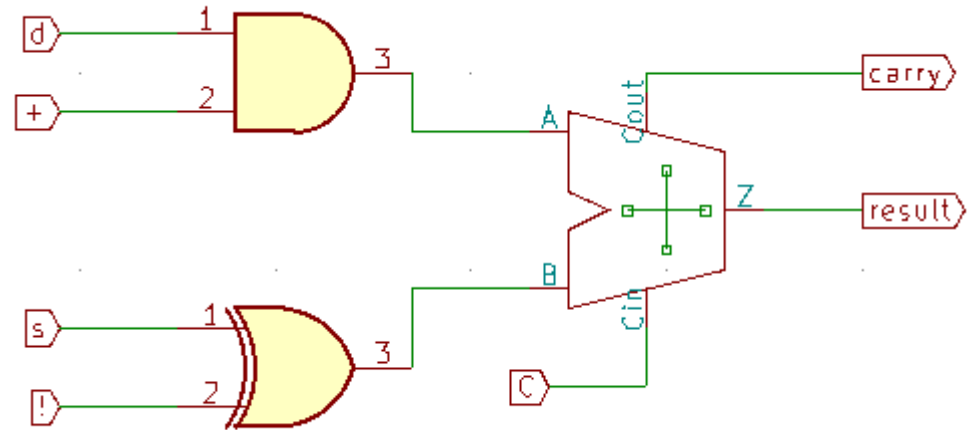
can be implemented with cascade as

```
r5 = {r3 << #2} + r1
```

Unlike the original code fragment, these two instructions don't destroy the value in r3. So though the architecture is a two address in general, it can have the functionality of three addresses (actually four) when needed. The second instruction's syntax is changed from "d op= s" to "d = {..} op s" and the cascaded instruction is placed in the curly brackets also in infix form. It is possible to have more than one level of cascade, but then the d field of intermediate instructions is wasted.

Math: op code is 0 0+!C

- + means that the destination is added to the source
- ! means that the source is bitwise inverted (one's complement)
- C means that the carry in is set

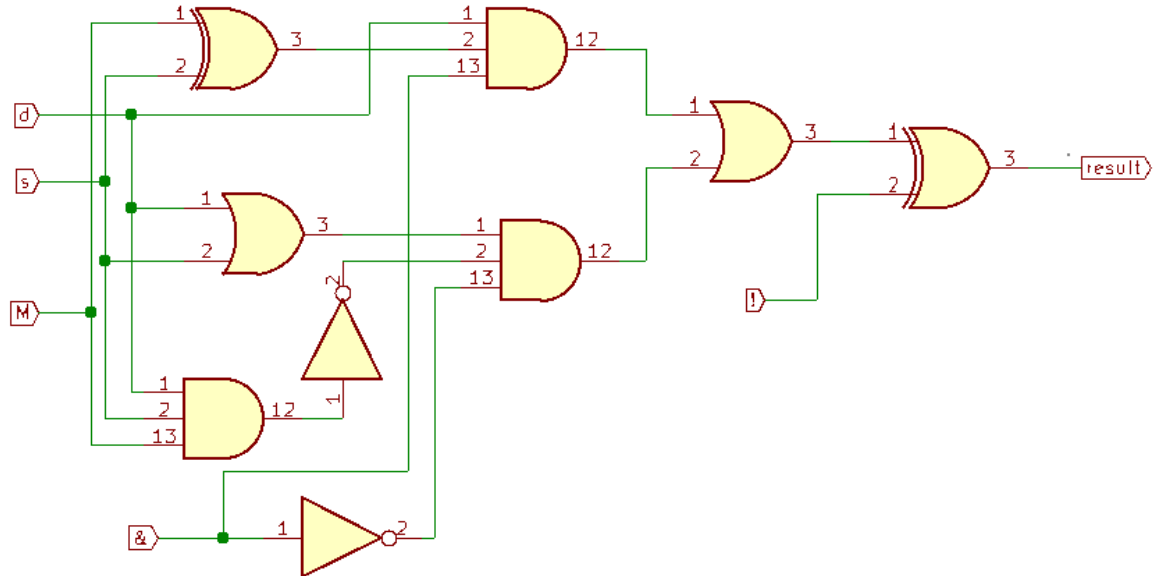


Math

		0?ds	1?ds	2?ds	3?ds
?0ds	move	$d = s$	$d = \#s$		
?1ds	increment	$d = s+1$	$d = \#s+1$		
?2ds	invert	$d = !s$	$d = !\#s$		
?3ds	negate	$d = \sim s$	$d = \sim \#s$		
?4ds	add	$d += s$	$d += \#s$	$\{d + s\}$	$\{d + \#s\}$
?5ds	add forcing carry	$d += s+1$	$d += \#s+1$	$\{d + s+1\}$	$\{d + \#s+1\}$
?6ds	subtract forcing borrow	$d -= s-1$	$d -= \#s-1$	$\{d - s-1\}$	$\{d - \#s-1\}$
?7ds	subtract	$d -= s$	$d -= \#s$	$\{d - s\}$	$\{d - \#s\}$

Logic: op code is 0 1&!M

- & means an AND operation, otherwise it is an OR
- ! means that output is bitwise inverted (one's complement)
- M means a modified version of the instruction



Logic (the 4 bits are !s!d s!d !sd sd)

		0?ds	1?ds	2?ds	3?ds
?8ds	0111 or	$d \mid= s$	$d \mid= \#s$	$\{d \mid s\}$	$\{d \mid \#s\}$
?9ds	0110 exclusive or	$d \wedge= s$	$d \wedge= \#s$	$\{d \wedge s\}$	$\{d \wedge \#s\}$
?Ads	1000 nor	$d \nmid= s$	$d \nmid= \#s$	$\{d \nmid s\}$	$\{d \nmid \#s\}$
?Bds	1001 equivalence	$d \nmid \wedge= s$	$d \nmid \wedge= \#s$	$\{d \nmid \wedge s\}$	$\{d \nmid \wedge \#s\}$
?Cds	0001 and	$d \&= s$	$d \&= \#s$	$\{d \& s\}$	$\{d \& \#s\}$
?Dds	0010 and invert	$d \&!= s$	$d \&!= \#s$	$\{d \&! s\}$	$\{d \&! \#s\}$
?Eds	1110 nand	$d \nmid \&= s$	$d \nmid \&= \#s$	$\{d \nmid \& s\}$	$\{d \nmid \& \#s\}$
?Fds	1101 nand invert	$d \nmid \&!= s$	$d \nmid \&!= \#s$	$\{d \nmid \&! s\}$	$\{d \nmid \&! \#s\}$

Logic: remaining rules

0000 clear	10d0	$d = \#0$
0011 destination	00dd	$d = d$
0100 other and invert	2Edd 0Cds	$d = \{d \text{ !\& } d\} \& s$
0101 source	00ds	$d = s$
1010 not source	02ds	$d = !s$
1011 other nand invert	2Edd 0Eds	$d = \{d \text{ !\& } d\} \text{ !\& } s$
1100 not destination	02dd	$d = !d$
1111 set	10dF	$d = \#-1$

Shift: op code is 1 00\$<

		4?ds	5?ds	6?ds	7?ds
?0ds	shift right	d >>= s	d >>= #s	{d >> s}	{d >> #s}
?1ds	shift left	d <<= s	d <<= #s	{d << s}	{d << #s}
?2ds	signed shift right	d \$>>= s	d \$>>= #s	{d \$>> s}	{d \$>> #s}
?3ds	rotate left	d <>= s	d <>= #s	{d <> s}	{d <> #s}

$a = "<" ? s : 32 - s$
 $x = ("\$" \& !"<") \& 32\{d[31]\}$
 $z[63:0] = \{x, d\} << a$
 $result = (z[31:0] \& "<") \mid (z[63:32] \& !(!"\$" \& "<"))$

Barrel shifter: $x[63:0] << a[4:0]$
 $z0[63:0] = a[0] ? \{x[62:0], 0\} : x$
 $z1[63:0] = a[1] ? \{z0[61:0], 2\{0\}\} : z0$
 $z2[63:0] = a[2] ? \{z1[59:0], 4\{0\}\} : z1$
 $z3[63:0] = a[3] ? \{z2[55:0], 8\{0\}\} : z2$
 $z[63:0] = a[4] ? \{z3[47:0], 16\{0\}\} : z3$

Multiply (optional)

		4?ds	5?ds	6?ds	7?ds
?4ds	multiply high	d *= s	d *= #s	{d * s}	{d * #s}
?5ds	multiply low	d *>= s	d *>= #s	{d *> s}	{d *> #s}
?6ds	multiply signed high	d \$*= s	d \$*= #s	{d \$* s}	{d \$* #s}
?7ds					

unsigned

signed

4 bit examples
with 8 bit
result:

hex (decimal)

4 (4)
x F (15)

3C (60)

C (12)
x F (15)

B4 (180)

4 (4)
x F (-1)

FC (-4)

C (-4)
x F (-1)

04 (4)

Comparison (calculation)

Greater than
 $d-s ==> !z \& c$

Equal
 $d-s ==> z$

Greater than or equal
 $d-s ==> c$

Not equal
 $d-s ==> !z$

Signed greater than
 $d-s ==> !z \& (n!^v)$

Carry on add
 $d+s ==> c$

Signed greater than or equal
 $d-s ==> n!^v$

$c = \text{result}[32]$
 $z = !(\backslash \text{result})$
 $n = \text{result}[31]$
 $v = (n \& !d[31] \& !s[31]) \mid (!n \& d[31] \& s[31])$

Comparison

		4?ds	5?ds	6?ds	7?ds
?8ds	Greater than	$d > s ? \dots$	$d > \#s ? \dots$	$\{d > s\}$	$\{d > \#s\}$
?9ds	Greater or equal	$d \geq s ? \dots$	$d \geq \#s ? \dots$	$\{d \geq s\}$	$\{d \geq \#s\}$
?Ads	Signed greater than	$d \$> s ? \dots$	$d \$> \#s ? \dots$	$\{d \$> s\}$	$\{d \$> \#s\}$
?Bds	Signed greater or equal	$d \$\geq s ? \dots$	$d \$\geq \#s ? \dots$	$\{d \$\geq s\}$	$\{d \$\geq \#s\}$
?Cds	Equal	$d == s ? \dots$	$d == \#s ? \dots$	$\{d == s\}$	$\{d == \#s\}$
?Dds	Not equal	$d != s ? \dots$	$d != \#s ? \dots$	$\{d != s\}$	$\{d != \#s\}$
?Eds	Carry on add	$d ++ s ? \dots$	$d ++ \#s ? \dots$	$\{d ++ s\}$	$\{d ++ \#s\}$
?Fds					

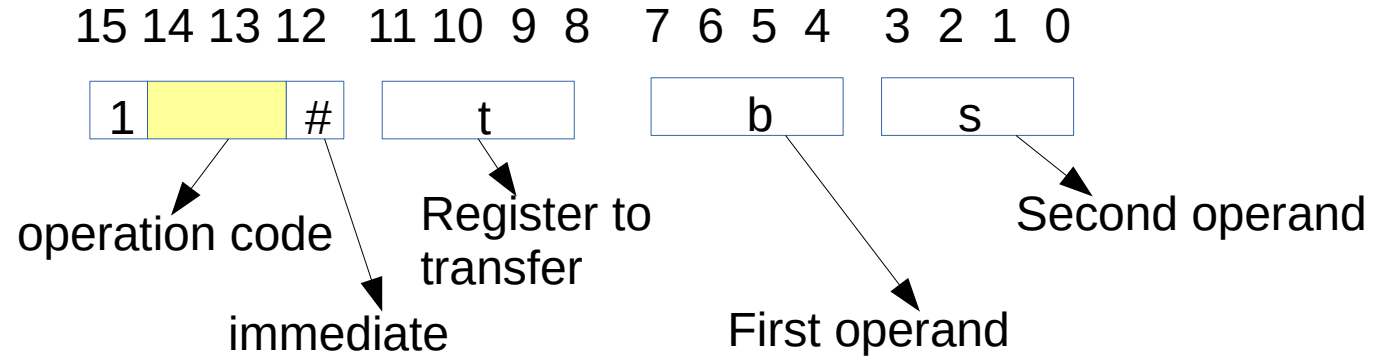
Comparison (trivia)

The comparison instructions will produce a 0 or a 1 result for the next instruction when cascade is used or will skip the next instruction in the normal case when the comparison fails (note that d is not modified)

Note that a comparison that wants less than results can use greater than and swap the operands. This doesn't work if the second operand is an immediate value, but since "if (r3<200) x1 else x2" does the same thing as "if (r3>=200) x2 else x1" it is normally not a problem.

The "d >= s" instruction can be used as a "borrow on subtract" to complement the "d ++ s" (carry on add) instruction.

Data Storage Instructions



Store	8tbs	$b[s] = t$	9tbs	$b[\#s] = t$
Load	Atbs	$t = b[s]$	Btbs	$t = b[\#s]$
Jump and Link	Ctbs	$====> b[s]/t$	Dtbs	$====> b[\#s]/t$
	E???		F???	

Load/Store

The address of the memory operation will be the sum of b and s (unless b is 0, in this case it will be just s and the "b" can be omitted):

8t0s $[s] = t$
9t0s $[\#s] = t$
At0s $t = [s]$
Bt0s $t = [\#s]$

This allows access to global variables and i/o ports without having to use up a register for the base.

Jump and Link

In the case of JL the t register will hold the previous value of the PC (unless t is 0, in which case the old value of the PC is discarded. "/t" can be omitted):

C0bs ====> b[s]
D0bs ====> b[#s]

The new value of the PC will be the sum of b and s (unless b is 0, in this case it will be the sum of s and the previous value of the PC. The "b" can be omitted):

Ct0s ====> [s]/t
Dt0s ====> [#s]/t
C00s ====> [s]
D00s ====> [#s]

Syntactic sugar for the assembler allows "====> label" and "====> label/t" to be written in place of "====> [#label-.]" and "====> [#label-.]/t" respectively.

Bytes: Store byte instruction

```
        ; code to load a byte from memory
Loadb:
        ; r1 is the address
        ; r2 is the destination
        ; r3 is scratch
r2 = r1[#0]
r3 = {r1 & #3} << #3
r2 = {r2 >> r3} & #255
```

This isn't too bad. If we redefine load to rotate the word it read on a non aligned address then just two instructions would be needed:

```
        ; code to load a byte from memory
Loadb:
r2 = r1[#0] ; might be rotated
r2 &= #255
```

```
        ; code to store a byte to memory
Storeb:
        ; r1 is the address
        ; r2 is the byte to be stored
        ; r3 and r4 are scratch
r4 = r1[#0] ; previous word
r3 = {r1 & #3} << #3
r2 &= #255 ; make sure it is a byte
r5 = #255
r5 <<= r3
r4 &~= r5 ; clear destination byte
r4 = {r2 << r3} | r4
r1[#0] = r4
```

We could define the remaining data storage instruction to be a store byte:

```
        ; code to store a byte to memory
Storeb:
r1[#0] B= r2
```

Interrupts

The focus in this design is simplicity, even at a cost in performance. Nearly all functionality is implemented in the external interrupt controller. An “irq” signal goes to the processor and an “iack” signal goes back to the controller.

The controller has a visible register for each interrupt level it can handle plus an extra one for the non existing “level 0”. The value of these registers is the address for the corresponding handler routine.

When some interrupt is pending the controller raises irq until the processor executes an iack cycle. The cycle looks like a memory read where the address is the value of the next PC (which is stored in the level 0 register or ignored if inside a handler) and the data read is the value to be loaded into the PC. Interrupts do no nest so even if a new, higher level request arrives it won't do anything until the current handler returns.

Return from Interrupt

After an iack cycle starts the execution of an interrupt handler, all further interrupt requests are ignored until the processor tries to read from the level 0 register. If there are any pending interrupt requests then a new irq/iack cycle happens but the value of the address bus is not stored in the level 0 register. That makes the next handler immediately follow the now finished one.

The simplest way to read L0 is to jump to it. That is a not very efficient “return from interrupt” instruction. The data returned from a normal read from L0 (not an iack cycle) causes a one instruction loop ($\text{==>}[\# - 4]$). The iack cycle will break out of the loop.

If no interrupts are pending then the value of L0 is returned as the data for the iack cycle, causing the processor to continue executing the user program.

Note that while the interrupt controller saves the user's PC, any other registers must be explicitly saved and then restored before the return.

Assembly Language Expressions

Numeric literals are always integers represented by a sequence of digits, optionally separated by underscores to make them more readable. The size depends on the context in which the expression is used. Though there is a default radix, any number literal can specify its representation by starting with the value of the highest digit in the base (radix), followed by a “\$” character and then followed by the actual digits.

We can have binary numbers like 1\$1101_1001, decimal numbers like 9\$42, hexadecimal numbers like F\$4AC0 or even octal numbers like 7\$1776. The digits after “9” are “A” to “Z” (lower case is equivalent to upper case).

Defined symbols are treated as their values.

The operations are, in order of highest to lowest precedence: *, /, % then +, - then <<, >>, \$>>, <> then & then ^ then |. Parenthesis override the precedence and must be used inside cascades so the top level operator is the instruction.

Assembly language directives: define

Symbols can be defined by

name: expression

A period character in an expression indicates the value of the PC for that instruction. An empty expression is equivalent to just a period, so is equivalent of defining a label in other assemblers. This means that labels must be on a line of their own instead of coming before an instruction.

If a symbol is defined more than once, only the last definition is used in the whole program and a warning is issued.

Registers in an assembly instruction are just expressions that evaluate from 0 to 15 (0 to 7 in Baby042) so besides the default names r0 to r15 more meaningful names can be defined.

Assembly language directives: configuration

Symbol names starting with “\$” are an error, unless they are in the list below and set to one of the indicated values. In that case they change the configuration of the assembler and/or the simulator.

Symbol	Valid Values	Default	Description
\$processor	42, 22, 18	42	Baby42, Baby22 or Baby042 (9\$18=7\$42)
\$radix_dump	1\$1 to Z\$Z	F\$F	Base of numbers in addresses and data
\$radix_source	1\$1 to Z\$Z	9\$9	Base of numbers in the source code
\$vresolution	0, 192, 384, 768	192	No video, 256x192, 512x386 or 1024x768
\$colorbits	1, 6, 12, 24	6	B&W, EGA, NeXT color, True Color

Assembly language directives: configuration (continued)

Symbol	Valid Values	Default	Description
\$framebuffer	Multiple of \$fboffset	F\$FFFF_4000	Where the video bitmap starts
\$fboffset	8KB, 16KB, 32KB, 64KB, 128KB, 256KB, 512KB, 1MB	16KB	Power of 2 larger than 1/3 of the bitmap size
\$ircontroller	Word aligned	F\$FFFF_FF00	Where the interrupt level registers are

Assembly language directives: include file

A line can be replaced with the contents of another file with:

```
< "file name"
```

As the assembler generates a single output and doesn't use a linker, the only way to develop with multiple files is to include the others into a top file. The tools can optionally show the contents of the file right below the directive adding "> " to each line. With multiple levels of include you can easily have lines starting with "> > > " or similar.

Assembly language directives: origin

The value of the PC can be changed with

`% expression`

This is equivalent of "org" in other assemblers. If 120 bytes need to be allocated, then "`% .+120`" will do the job.

Instructions should always be aligned on an even byte. This expression can force that to be the case:

`% ((.+1)/2)*2`

Assembly language directives: constants

Constants can be inserted into the instruction stream with

`# expr1, expr2, expr3, B# expr4, expr5, H# expr6, # expr7, expr8`

The `#` interprets the following expressions as 32 bit values, while `B#` inserts 8 bit values and `H#` 16 bit values.

Placing ASCII characters between two `"` has the same effect as `B#` followed by the list of the characters' numeric value.

Example 1

Jan Gray used this simple C program to illustrate the assembly language of his XR16 processor (<http://www.fpgacpu.org/papers/xsoc-series-drafts.pdf>)

```
typedef struct TN {
    int k;
    struct TN *left, *right;
} *T;

T search(int key, T p) {
    while (p && p->k != key)
        if (p->k < key)
            p = p->right;
        else
            p = p->left;
    return p;
}
```


; Example 1 Baby42 machine language and assembly with LCC style

k: 0

left: 4

right: 8

search: ; r3=key r4=p r9=scratch r2=return

D008 000E ====> L3

L2:

B940 r9 = r4[#k]

4893 D006 r9 > r3 ? ====> L5

B448 0008 r4 = r4[#right]

D002 ====> L6

L5:

B444 r4 = r4[#left]

L6:

L3:

0094 r9 = r4

5C90 D008 0008 r9 == #0 ? ====> L7

B940 r9 = r4[#k]

4D93 D008 FFE2 r9 != r3 ? ====> L2

L7:

0024 r2 = r4

L1:

D0D0 ====> r13[#0] ; return address was in r13

```

; Example 1 Baby42 machine language and assembly with better compiler
;      that optimizes jumps to jumps and jumps to returns
k:      0
left:   4
right:  8

```

```

search: ; r3=key r4=p r9=scratch r4=return value
5C90 D0D0      r4 == #0 ? ====> r13[#0] ; return
B940          r9 = r4[#k]
5C93 D0D0      r9 == r3 ? ====> r13[#0]
4893 D008 0008  r3 > r9 ? ====> L5
B448 0008      r4 = r4[#right]
D008 FFE8      ====> search

L5:
B444          r4 = r4[#left]
D008 FFE2      ====> search

```

Example 2

Here is the famous Sieve of Eratosthenes benchmark published in Byte magazine:

```
1 SIZE = 8190
2 DIM FLAGS(8191)
3 PRINT "Only 1 iteration"
5 COUNT = 0
6 FOR I = 0 TO SIZE
7   FLAGS (I) = 1
8 NEXT I
9 FOR I = 0 TO SIZE
10  IF FLAGS (I) = 0 THEN 18
11  PRIME = I+I + 3
12  K = I + PRIME
13  IF K > SIZE THEN 17
14  FLAGS (K) = 0
15  K = K + PRIME
16 GOTO 13
17 COUNT = COUNT + 1
18 NEXT I
19 PRINT COUNT," PRIMES"
```

```

size: 8190
true: r10
false: r11
flags: r12
flg:
???? .. ??    % .+8191
count: r1
i:    r2
prime: r3
k:    r4
text1:
    4F 6E6C    "Only 1 iteration", B# 0
7920 3120
6974 6572
6174 696F
6E00
text2:
2050 5249    " PRIMES", B# 0
4D45 5300
sieve:
10C0    flags = #flg
10A1    true = #1
10B0    false = #0
1008 1FFF    r0 = #text1

```

```

DD08 ????    ====> PrintText/r13
                ; library subroutine expects
                ; return address in r13
1010    count = #0
1020    i = #0
for1:
5828 1FFE    i > #size ? ====> endfor1
D008 0008
8AC2        flags[i] = true
1421        i += #1
D008 FFF0    ====> for1
endfor1:
1020        i = #0
for2:
5828 1FFE    i > #size ? ====> endfor2
D008 0028
A5A2        r5 = flags[i]
5C50 D008    r5 == #0 ? ====> L18
001A
2422 1433    prime = {i+i} + #3
2423 1440    k = {i + prime} + #0

```

```

L13:
5848 1FFE      k > #size ? ====> L17
D008 0008
8BA4          flags[k] = false
0443          k += prime
D008 FFF0      ====> L13

L17:
1411          count += #1

L18:
1421          i += #1
D008 FFDO      ====> for2
endfor2
0001          r0 = count
DD08 ????      ====> PrintInt/r13
1008 2010      r0 = #text2
DD08 ????      ====> PrintText/r13
D008 ????      ====> Stop

```

There is also a C version of the benchmark in the Wikipedia article (https://en.wikipedia.org/wiki/Byte_Sieve). But Basic, specially a version as primitive as this, corresponds a lot more directly to the assembly language implementation.

Normally code that calls other subroutines should save the return address on the stack, but in this case we end with a jump instead of a return.

Since the address of flags happens to be 0 (it is a byte pointer), it would be possible to use "i[#FLG]" in the load and store instructions instead of "flag[i]", which would save one register. But that is not something that would normally happen, so it wasn't used in this example either.

Example 3: ARM and Baby42

```
AREA LOG, CODE, READONLY
EXPORT log
; r0 = input variable n
; r0 = output variable m (0 by default)
; r1 = output variable k ( $n \leq 2^k$ )
```

Log

```
MOV r2, #0 ; set m = 0
```

```
MOV r1, #-1 ; set k = -1
```

log_loop

```
TST r0, #1 ; test LSB(n) == 1
```

```
ADDNE r2, r2, #1 ; set m = m+1 if true
```

```
ADD r1, r1, #1 ; set k = k+1
```

```
MOVS r0, r0, LSR #1 ; set n =  $n \gg 1$ 
```

```
BNE log_loop ; continue if  $n \neq 0$ 
```

```
CMP r2, #1 ; test m == 1
```

```
MOVEQ r0, #1 ; set m = 1 if true
```

log_rtn

```
MOV pc, lr
```

END

40 bytes, hence Thumb

```
; AREA LOG, CODE, READONLY
```

```
; EXPORT log
```

```
; r0 = input variable n
```

```
; r0 = output variable m (0 by default)
```

```
; r1 = output variable k ( $n \leq 2^k$ )
```

Log:

```
r2 = #0 ; set m = 0
```

```
r1 = #-1 ; set k = -1
```

log_loop:

```
{r0 & #1} == #1 ? r2 += #1 ; test LSB(n) == 1
```

```
; set m = m+1 if true
```

```
r1 += #1 ; set k = k+1
```

```
r0 >>= #1 ; set n =  $n \gg 1$ 
```

```
r0 != #0 ? =====>log_loop ; continue if  $n \neq 0$ 
```

```
r2 == #1 ? r0 = #1 ; test m == 1
```

```
; set m = 1 if true
```

log_rtn:

```
=====>[lr]
```

26 bytes

What is missing?

- No floating point math
- No exceptions (if the missing op codes are defined as either NOPs or redundant)
- No MMU (memory management unit) or supervisor mode
- No debug interface (JTAG or control panel)

Variation 1: baby22

- It is possible to have a 16 bit datapath so the processor would have both 2 byte instructions and 2 byte data. This is limited to 64KB of memory, but that is enough for many applications
- The only change from baby42 is that immediate value -7 does not indicate a 32 bit immediate

Variation 2: baby042

- Operating on groups of 6 bits instead of 8 bit bytes we could have 24 bit data and 12 bit instructions. It is more natural to use octal instead of hexadecimal numbers
- With 4 fewer bits in each instruction it would be better to have 8 registers instead of 16. Instead of a dedicated immediate bit, when the second operand is register 6 or 7 the actual value is a 24 or 12 bit immediate after the instruction.
- 0XDS: $d=s$, $d=s+1$, $d=!s$, $d=\sim s$,
- 1XDS: $d+=s$, $\{d+s\}$, $d+=s+1$, $\{d+s+s\}$, $d-=s-1$, $\{d-s-1\}$, $d-=s$, $\{d-s\}$
- 2XDS: $d|=s$, $\{d|s\}$, $d^=s$, $\{d^s\}$, $d\&s$, $\{d\&s\}$, $d!\&s$, $\{d!\&s\}$
- 3XDS: $d>>=s$, $\{d>>s\}$, $d<<=s$, $\{d<<s\}$, $d\$>>=s$, $\{d\$>>s\}$, $d<<=s$, $\{d<<=s\}$
- 4XDS: $d>s?$, $\{d>s\}$, $d>=s?$, $\{d>=s\}$, $d==s?$, $\{d==s\}$, $d!=s?$, $\{d!=s\}$
- 5TBS: $===> b[s]/t$
- 6TBS: $b[s] = t$
- 7TBS: $t = b[s]$